

Design a Payment System

Multi-step Processes

By Evan King · Updated Jul 11, 2025 · hard · Asked at:



Try This Problem Yourself

Practice with guided hints and real-time feedback



Start Practice

The video player displays a presentation slide titled "Payment System like Stripe". The slide content reads: "Payment processing systems like Stripe allow business (referred to throughout as merchants) to accept payment from customers, without having to build their own payment processing infrastructure. Customer input their payment details on the merchant's website, and the merchant sends the payment details to Stripe. Stripe then processes the payment and returns the result to the merchant." The video player interface includes a progress bar at 0:00 / 1:03:59, a "Premium Content" label, and a small video inset of a man in the bottom right corner.

Understanding the Problem

What is Stripe?

Payment processing systems like Stripe allow business (referred to throughout this breakdown as merchants) to accept payment from customers, without having to build their own payment processing infrastructure. Customer input their payment details on the merchant's website, and the merchant sends the payment details to Stripe. Stripe then processes the payment and returns the result to the merchant.

Functional Requirements

Core Requirements

1. Merchants should be able to initiate payment requests (charge a customer for a specific amount).
2. Users should be able to pay for products with credit/debit cards.
3. Merchants should be able to view status updates for payments (e.g., pending, success, failed).

Below the line (out of scope):

- Customers should be able to save payment methods for future use.
- Merchants should be able to issue full or partial refunds.
- Merchants should be able to view transaction history and reports.
- Support for alternative payment methods (e.g., bank transfers, digital wallets).
- Handling recurring payments (subscriptions).
- Payouts to merchants.

Non-Functional Requirements



Before defining your non-functional requirements in an interview, it's wise to inquire about the scale of the system as this will have a meaningful impact on your design. In this case, we'll be looking at a system handling about 10,000 transactions per second (TPS) at peak load.

Core Requirements

1. The system should be highly secure
2. The system should guarantee durability and auditability with no transaction data ever being lost, even in case of failures.
3. The system should guarantee transaction safety and financial integrity despite the inherently asynchronous nature of external payment networks
4. The system should be scalable to handle high transaction volume (10,000+ TPS) and potentially bursty traffic patterns (e.g., holiday sales).

Below the line (out of scope):

- The system should adhere to various financial regulations globally (depending on supported regions).
- The system should be extensible to easily add new payment methods or features later.

Here's how it might look on your whiteboard:

Functional Requirements

- Merchants should be able to create payments
- Users should be able to pay via credit/debit card
- Merchants should be able to view status updates for payments

Non-functional Requirements

- The system should be highly secure
- Transaction safety and financial integrity (no double charge, etc)
- Strong durability (can't lose transactions)
- Scale to 10k TPS



The Set Up

Defining the Core Entities

Let's start by identifying the core entities we'll need. I prefer to begin with a high-level overview before diving into specifics as this helps establish the foundation we'll build upon. That said, if you're someone who finds value in outlining the complete schema upfront with all columns and relationships defined, that's perfectly fine too! There's no single "right" approach in an interview, just do what works best for you. The key is to have a clear understanding of the main building blocks we'll need.

To satisfy our key functional requirements, we'll need the following entities:

1. **Merchant:** This entity will store information about the businesses using our payment platform, including their identity details, bank account information, and API keys.
2. **PaymentIntent:** This represents the merchant's intention to collect a specific amount from a customer and tracks the overall payment lifecycle from creation to completion. It owns the state machine from created → authorized → captured / canceled / refunded, and enforces idempotency for retries.
3. **Transaction:** This represents a polymorphic money-movement record linked back to one PaymentIntent. Types include Charge (funds in), Refund (funds out), Dispute (potential reversal), and Payout (merchant withdrawal). Each row carries amount, currency, status, timestamps, and references to both the intent and the merchant. For simplicity, we'll only be focused on Charges as everything else is out of scope.

It's important to clarify the distinction between PaymentIntent and Transaction at this point as this can easily cause confusion. The relationship between them is one-to-many: a single PaymentIntent can have multiple Transactions associated with it. For example:

- If a payment attempt fails due to insufficient funds, the merchant might retry with the same PaymentIntent ID but a different Transaction will be created.
- For partial payments, multiple Transactions might be linked to the same PaymentIntent.
- For refunds, a new Transaction with a negative amount would be created and linked to the original PaymentIntent.

From the merchant's perspective, they simply create a PaymentIntent and our system handles all the transaction complexities internally, providing a simplified view of the overall payment status.

In the actual interview, this can be as simple as a short list like this. Just make sure you talk through the entities with your interviewer to ensure you are on the same page.



A **Transaction** is a bit of a simplification to ensure we can focus on the important parts of the design with our limited time. In production you'd likely break this out into discrete types— **Charge** , **Refund** , **Dispute** , **Payout** , and the double-entry **LedgerEntry** rows that actually move balances. For interview purposes, though, collapsing them under a single polymorphic **Transaction** keeps the mental model tight: one **PaymentIntent** can spawn many money-movement events, each stamped with an amount, currency, status, and timestamps. That level of detail is enough to reason about idempotency, auditability, and failure handling without drowning in implementation minutiae.

Core Entities

- Merchant
- PaymentIntent
- Transaction ➕

API or System Interface

The API is the main way merchants will interact with our payment system. Defining it early helps us structure the rest of our design. We'll start simple and, as always, we can add more detail as we go. I'll just create one endpoint for each of our core requirements.

First, merchants need to initiate `PaymentIntent` requests. This will happen when a customer reaches the checkout page of a merchant's website. We'll use a `POST` request to create the `PaymentIntent`.

```
POST /payment-intents -> paymentIntentId
{
  "amountInCents": 2499,
  "currency": "usd",
  "description": "Order #1234",
}
```

Next, the system needs to securely accept and process payments. Since we're focusing on credit/debit cards initially, we'll need an endpoint for that:

```
POST /payment-intents/{paymentIntentId}/transactions
{
  "type": "charge",
  "card": {
    "number": "4242424242424242",
    "exp_month": 12,
    "exp_year": 2025,
    "cvc": "123"
  }
}
```



In a real implementation, we'd never pass raw card details directly to our backend like this. We'd use a secure tokenization process to protect sensitive data. We'll get into the details of how we handle this securely when we get further into our design. In your interview, I would just callout that you understand this data will need to be encrypted.

Finally, merchants need to check the status of payments. This can be done with a simple `GET` request like so:

```
GET /payment-intents/{paymentIntentId} -> PaymentIntent
```

The response would include the payment's current status (e.g., "pending", "succeeded", "failed") and any relevant details like error messages for failed payments.

For a more real-time approach (and one that actually mirrors the industry standard), we could also provide webhooks that notify merchants when payment statuses change. In this way, the merchant would provide us with a callback URL that we would `POST` to when the payment status changes, allowing them to get real-time updates on the status of their payments.

the status of their payment:

```
POST {merchant_webhook_url}
{
  "type": "payment.succeeded",
  "data": {
    "paymentId": "pay_123",
    "amountInCents": 2499,
    "currency": "usd",
    "status": "succeeded"
  }
}
```



Designing a webhook callback system is often a standalone question in system design interviews. Designing a payment system, complete with webhooks, would be a lot to complete in the time allotted. Have a discussing with your interviewer early on so you're on the same page about what you're building.


High-Level Design

For our high-level design, we're simply going to work one-by-one through our functional requirements.

1) Merchants should be able to initiate payment requests

When a merchant wants to charge a customer, they need to initiate a payment request. In our system, this is accomplished through the creation of a `PaymentIntent`. As we stated in our Core Entities, a `PaymentIntent` represents the merchant's intention to collect a specific amount from a customer and tracks the lifecycle of the payment from initiation to completion.

Let's start by laying out the core components needed for this functionality:



- API Gateway:** This serves as the entry point for all merchant requests. It handles authentication, rate limiting, and routes requests to the appropriate microservices.
- Payment Service:** This microservice is responsible for creating and managing `PaymentIntents`. It interfaces with our database to store payment information.
- Database:** A central database that stores all system data including `PaymentIntents` records (with their current status, amount, currency, and associated metadata) and merchant information (API keys, business details, and configuration preferences).

Here's the flow when a merchant initiates a payment:

Here's how this flow when a merchant initiates a payment:

1. The merchant makes an API call to our system by sending a POST request to `/payment-intents` with details like amount, currency, and description.
2. Once authenticated, the request is routed to the PaymentIntent Service (more on this later).
3. The PaymentIntent Service creates a new PaymentIntent record with an initial status of "created" and stores it in the Database.
4. The system generates a unique identifier for this PaymentIntent and returns it to the merchant in the API response.

This PaymentIntent ID is crucial as it will be used in subsequent steps of the payment flow. The merchant will typically embed this ID in their checkout page or pass it to their client-side code where it will be used when collecting the customer's payment details.

At this stage, no actual charging has occurred. We've simply recorded the merchant's intention to collect a payment and created a reference that will be used to track this payment throughout its lifecycle. The PaymentIntent exists in a "created" state, awaiting the next step where payment details will be provided and processing will begin.

2) Users should be able to pay for products with credit/debit cards.

Now that we have a PaymentIntent created, the next step is to securely collect payment details from the customer and process the payment. It's important to understand that payment processing is inherently asynchronous - the initial authorization response from the payment network is just the first step in a longer process that can take minutes or even days to fully complete. This is because payment networks need time to handle things like fraud checks, chargeback requests, etc.


Pattern: Multi-step Processes


Payment processing is a perfect example of the multi-step processes pattern. A single payment goes through multiple stages: authorization, capture, settlement, and potentially refunds or disputes. Each step can fail independently and may require retries or compensation. This requires careful state management and orchestration to ensure the payment workflow completes successfully while handling failures gracefully.

[Learn This Pattern](#)

Our system needs to handle this asynchronous nature by maintaining the state of each payment and transaction, and keeping merchants informed of status changes throughout the entire lifecycle.

Let's expand our architecture to handle this critical part of the payment flow:





Process Payment Flow



1. **Transaction Service:** A dedicated microservice responsible for receiving card details from the merchant server, managing transaction records throughout the payment lifecycle, and interfacing directly with external payment networks like Visa, Mastercard, and banking systems.
2. **External Payment Network:** Shown as a dotted line in our diagram to demonstrate that this is external to our actual system, though a crucial part of the payment flow. These are the payment networks (Visa, Mastercard, etc.) and banking systems that actually authorize and process the financial transactions.

Here's how the flow works when a customer enters their payment details:

1. The customer enters their credit card information into a payment form on the merchant's website.
2. The merchant collects this data and sends it to our Transaction Service along with the original PaymentIntent ID.
3. The Transaction Service creates a transaction record with status "pending" in our system.
4. The Transaction Service directly handles the payment network interaction: a. Connects to the appropriate payment network and sends the authorization request b. Receives the initial response (approval/decline) c. Updates the transaction record with the initial status d. Continues to listen for callbacks from the payment network over the secure private connection e. When additional status changes occur (settlement, chargeback, etc.), receives callbacks and updates records accordingly
5. The Transaction Service updates the PaymentIntent status as the transaction progresses through its lifecycle.



While it would be highly unusual for any of these details to be covered in anything but the most specialized interviews, you may be wondering how connections to payment networks work.


In short, payment networks operate on private, highly secure networks that are completely separate from the public internet. To connect with these networks, payment processors must establish dedicated connections through specialized hardware security modules (HSMs) and secure data centers that meet stringent [Payment Card Industry Data Security Standard \(PCI DSS\)](#) requirements. These private networks use proprietary protocols with specific message formats (such as ISO 8583) and require formal certification processes to gain access. Unlike typical [REST APIs](#), these connections often involve binary protocols, leased lines, and VPN tunnels with mutual TLS authentication.

This simplified architecture keeps the Transaction Service as the single point of responsibility for both managing our internal transaction records and interfacing with external payment networks. While this combines multiple concerns in one service, it reduces complexity and eliminates unnecessary network hops while still maintaining security through proper PCI compliance within the service.

3) The system should provide status updates for payments

After a payment is initiated and processed, merchants need a reliable way to determine its current status. This information is very important for business operations! Merchants need to know when a payment succeeds to trigger fulfillment actions like shipping physical products, granting access to digital content, or confirming reservations. Likewise, they need to know when payments fail so they can notify customers or attempt alternative payment methods.

Let's see how our existing architecture supports this functionality:



Since we already have a PaymentIntent Service that manages PaymentIntents, we can leverage this same service to provide status updates to merchants. There's no need to create a separate service just for checking statuses.

Here's how the flow works when a merchant checks a payment's status:

1. The merchant makes a GET request to `/payment-intents/{paymentIntentId}` to retrieve the current status of a specific PaymentIntent.
2. The API Gateway validates the merchant's authenticity and routes the request to the PaymentIntent Service.
3. The PaymentIntent Service queries the database for the current state of the PaymentIntent, including its status, any error messages (if failed), and related transaction details.
4. The service returns this information to the merchant in a structured response format.

This simple polling mechanism allows merchants to programmatically check payment statuses and integrate the results into their business workflows.

The PaymentIntent can have various statuses throughout its lifecycle, such as:

- **created** : Initial state after the merchant creates the PaymentIntent
- **processing** : PaymentIntent details received and being processed
- **succeeded** : PaymentIntent successfully processed
- **failed** : Payment processing failed (with reason)

While this polling approach works well for many use cases, it's not ideal for real-time updates or high-frequency status checks. In a deep dive later, we'll explore how webhooks can be implemented to provide push-based notifications that eliminate the need for polling and reduce latency between payment completion and fulfillment actions.

Now that we've covered all three functional requirements, we have a basic payment processing system that can create payments, process payments securely, and provide status updates to merchants.

Potential Deep Dives

At this point, we have a basic system that satisfies the core functional requirements of our payment processing system. Merchants can initiate payments by creating a PaymentIntent, customers can pay with credit/debit cards

system. merchants can initiate payments by creating a payment intent, customers can pay with credit/debit cards, and merchants can view payment status updates. However, our current design has significant limitations, particularly around transaction safety, durability, and scaling to handle high volumes. Let's look back at our non-functional requirements and explore how we can improve our system to handle 10,000+ TPS with strong consistency and guaranteed durability.

1) The system should be highly secure

Let's start with security. For a payment processing system, there are two main things we care about when it comes to guaranteeing the security of the system.

1. Is the person/merchant making the payment request who they say they are?
2. Are we protecting customer personal information so that it can't be stolen or compromised?


Starting with #1, we need to validate that merchants connecting to our system are who they claim to be. After all, we're giving them the ability to charge people money, we better make sure they're legit! Most payment systems solve this with API keys, but there are different approaches with varying levels of security. Here are some options.

Good Solution: Basic API Key Authentication

Approach

We can use standard, static API keys as the primary authentication mechanism for merchants. When a merchant onboards to our payment platform, we generate a unique API key (typically a random string like `pk_live_51NzQRtGswQnXYZ8o`) and store it in our database associated with the merchant's account.

For each API request, merchants include this key in the request headers, typically as `Authorization: Bearer {api_key}` or a custom header like `X-API-Key: {api_key}`. When our API Gateway receives a request, it extracts the API key, looks it up in the database, and identifies the corresponding merchant. If the key is valid, the request is authenticated and processed.



Challenges

Since API keys are static credentials passed in every request, they're vulnerable to interception through network sniffing, especially if transmitted over insecure connections. Once intercepted, an attacker can replay these credentials indefinitely, as there's no mechanism to detect or prevent replay attacks.

Additionally, and even more likely, these keys often end up hardcoded in client applications or configuration files, increasing the risk of exposure through code repositories or system backups.

For a financial system processing sensitive payment data, these vulnerabilities represent unacceptable risks that could lead to fraudulent transactions, data breaches, and significant financial and reputational damage.

Great Solution: Enhanced API Key Management with Request Signing

Approach

To improve on the good solution, we can implement request signing which ensures that API requests are authentic, unmodified, and cannot be replayed by attackers. We can accomplish this with a combination of public API keys like before (to identify the merchant) and private secret keys (used to generate time bound signatures).

During merchant onboarding, we provide two keys: a public API key for identification and a private secret key stored securely on the merchant's server (never in client-side code). These keys are used for authenticating the merchant's server with our payment system, which is separate from how we handle customer card data (covered in the next section).

For each API request, the merchant's server generates a digital signature by hashing the request details (method, endpoint, parameters, body) along with a timestamp and a unique nonce using their secret key. This signature proves the request's authenticity and prevents tampering. This way, even if replayed, we'd know that the timestamp was outside our acceptable window or that the nonce was already used, allowing us to reject the request.

```
// Example request with signature
{
  "method": "POST",
  "path": "/payment-intents/{paymentIntentId}/transactions",
  "body": {
    // body here
  },
  "headers": {
    "Authorization": "Bearer pk_live_51NzQRtGswQnXYZ8o", // API Key
    "X-Request-Timestamp": "2023-10-15T14:22:31Z", // Timestamp
    "X-Request-Nonce": "a1b2c3d4-e5f6-7890-abcd-ef1234567890", // Nonce
    "X-Signature": "sha256=7f83b1657ff1fc53b92dc18148a1d65dfc2d4b1fa3d677284addd200126d9069" // H
  }
}
```

When our API Gateway receives a request, it:

1. Retrieves the merchant's secret key based on the provided API key
2. Recreates the HMAC signature using the same algorithm (SHA-256), secret key, and request data
3. Compares the calculated signature with the one provided in the request headers
4. Validates that the timestamp is within an acceptable time window (typically 5-15 minutes)

5. Ensures the nonce hasn't been used before within the valid time window by checking the cache/DB

This approach uses HMAC (Hash-based Message Authentication Code) with a shared secret key, ensuring both authenticity (the request came from the merchant with the valid secret) and integrity (the request data hasn't been tampered with).

Now let's look at #2 - protecting sensitive customer data throughout the payment process. Allowing a bad actor to get a hold of someone else's credit card information can lead to fraud, identity theft, and a total loss of customer trust. Plus, there are strict regulations like PCI DSS that mandate how payment data must be handled. Let's look at the different approaches to tackling this challenge.

Bad Solution: Server-Side Payment Data Collection ^

Approach

Merchants collect and process card data on their own servers before sending it to our payment system. The customer enters card details on the merchant website, these card details are forwarded to the merchant's server over HTTPS, where they encrypt and transmit the data to our payment service. The merchant's server then receives authorization responses and transaction confirmations from our system.

Challenges

A key security consideration is that sensitive card data should never touch the merchant's servers. This is not only a best practice but also a requirement for PCI DSS.

Every merchant server becomes a potential attack target. If any merchant has vulnerabilities, customer card data is at risk. Merchants also bear enormous liability; a single breach could result in devastating financial penalties and reputation damage.


Good Solution: Client-Side iFrame Isolation ^

Approach

Our big learning from the previous approach is that private data (like card numbers) should never hit the merchant's server. By handling this data ourselves, we protect merchants from compliance requirements and liability risks. Credit card information is transmitted directly from the client to our servers over HTTPS, bypassing the merchant's systems entirely.

To accomplish this, we provide a JavaScript SDK that loads iframes directly on the merchant's client. Our SDK creates a secure frame on the merchant's webpage that loads directly from our domain. When customers enter card information, it goes straight to our servers without touching the merchant's systems.

The iframe creates a security boundary—the browser's same-origin policy prevents the merchant's code from accessing data inside the frame.



You can learn more about iframes and how they work [here](#).

Challenges

While this approach prevents card data from reaching merchant servers, it still has security limitations. The iframe's security depends on browser security policies, which could be compromised by sophisticated attacks. Additionally, the card data is only protected by HTTPS during transmission, meaning it could be exposed if the iframe is compromised. The iframe approach also limits customization options and can cause usability issues on mobile browsers.

Great Solution: iFrame + Encryption

Approach


The most secure approach builds upon the iframe solution while addressing its limitations through multiple layers of security. Our payment system generates and manages encryption keypairs, and our JavaScript SDK utilizes our public key for encryption.

When a customer enters their card details in our iframe, our JavaScript SDK immediately encrypts the data using this public key before it even leaves their device. This means the card data never exists in an unencrypted form outside the customer's browser. Even if the merchant's site is compromised, the attacker would only get encrypted data that they cannot decrypt.

The encrypted data is then transmitted over HTTPS to our servers, where we use the corresponding private key (securely stored in Hardware Security Modules) to decrypt it. While HTTPS provides transport security, it's now protecting already-encrypted data, adding an extra layer of protection.

This multi-layered approach means that even if one security mechanism fails (e.g., if the iframe is compromised), the other layers continue to protect the card data. The data is encrypted before it leaves the customer's device, protected during transmission, and securely processed on our servers.





2) The system should guarantee durability and auditability with no transaction data ever being lost, even in case of failures.

For a payment system, the worst thing you can do is lose transaction data. It would be both a financial and legal disaster. Every transaction represents real money moving between accounts, and regulations like PCI-DSS, SOX compliance, and financial auditing standards require us to maintain complete, immutable records of every payment attempt, success, and failure.

We need to track not just what the current state is, but the entire sequence of events that led to that state. When a customer disputes a charge six months later, we must be able to prove exactly what happened: when the payment was initiated, what amount was authorized, when it was captured, and whether any refunds were processed. A single missing record could mean inability to defend against chargebacks, failed compliance audits, or worse—being unable to determine the true state of customer accounts.

This durability requirement intersects with several other system needs. As we'll explore later, webhook delivery requires knowing which events have been sent to which merchants. Reconciliation with payment networks demands comparing our records against theirs. Fraud detection needs to analyze patterns across transaction history. All of these capabilities depend on having a durable, queryable audit trail as their foundation.

Let's examine different approaches to achieving this durability and auditability:

Bad Solution: Single Database with Current State Only

Approach

The simplest (and worst!) approach is to store payment data in a traditional relational database, updating records in place as payment states change. When a payment moves from "pending" to "captured", we simply update the status field:

```
UPDATE payment_intents
SET status = 'captured',
    captured_at = NOW(),
    updated_at = NOW()
WHERE payment_id = 'pay_123';
```

This gives us the current state of every payment, which seems sufficient for basic operations. The database provides ACID guarantees, so we know updates are atomic and consistent. We can query current payment statuses quickly and build simple dashboards showing successful, failed payments.

statuses quickly and build simple dashboards showing successful vs. failed payments.

For durability, we rely on database backups—perhaps daily snapshots plus transaction logs that let us restore to any point in time if disaster strikes.

Challenges

The massive problem is that we've thrown away history. When the update above executes, we lose the fact that this payment was in "pending" state, when that transition happened, and who initiated it. If a customer claims they were incorrectly charged, we cannot prove the sequence of events that led to the charge.

More critically, this approach is vulnerable to application bugs. A faulty deployment that incorrectly updates payment statuses would silently corrupt our data. By the time we notice, the original state is gone forever. Even with database backups, we might not know which records were affected or what their correct values should be.

From a compliance perspective, this approach fails most financial audit requirements. Auditors need to see not just current balances but the complete transaction history. They need proof that records haven't been tampered with. A mutable database where history can be overwritten provides neither capability.

Lastly, when payment disputes arise, customer support must manually piece together what happened from server logs, if they still exist. Debugging production issues becomes archeology rather than investigation.

Good Solution: Separate Audit Tables

Approach

We can solve for most of the problem above with a more sophisticated approach that maintains separate audit tables alongside the main transaction tables. Every state change triggers an insert into an append-only audit log:

```
BEGIN TRANSACTION;
-- Update main table
UPDATE payments
SET status = 'captured',
    updated_at = NOW()
WHERE payment_id = 'pay_123';

-- Insert audit record
INSERT INTO payment_audit_log (
  payment_id,
  change_type,
  old_status,
  new_status,
  changed_by,
  changed_at,
  metadata
) VALUES (
  'pay_123',
  'status_change',
  'authorized',
  'captured',
  'payment_service',
  NOW(),
```



```
'{"amount": 2500, "auth_code": "ABC123"}'  
);  
COMMIT;
```

This preserves history while keeping the main tables optimized for queries. The audit tables are append-only—we never UPDATE or DELETE audit records, providing some immutability guarantees. For additional safety, we might restrict database permissions so audit tables are insert-only even at the database level.

We can now answer compliance questions by querying the audit log. Customer disputes can be investigated by examining the complete state transition history. If the main tables become corrupted, we could theoretically reconstruct the current state by replaying all audit records.

Challenges

While this is much better, it still has significant limitations. The main problem is that application code must remember to perform both operations—updating the main table AND inserting the audit record. If a developer forgets the audit insert, or a bug skips it, we silently lose our audit trail with no way to recover the missing history.

Additionally, this approach forces your high-volume operational database to also store years of audit history. The same database handling 10,000 TPS of real-time payment processing must also manage ever-growing audit tables that will eventually contain billions of historical records. This coupling makes it impossible to optimize each use case independently—audit queries for compliance slow down operational queries, and operational schema changes become constrained by audit requirements.

Great Solution: Database + Change Data Capture + Event Stream



Approach

A more robust approach, used by payment processors like Stripe, separates operational and audit concerns while guaranteeing consistency between them. We still use a traditional database for merchant-facing operations, but audit durability comes from an immutable event stream populated via Change Data Capture (CDC).

Here's how it works:

1. **Operational Database:** Handles merchant API requests with optimized tables for current state queries. No audit tables needed—just pure operational data models.
2. **Change Data Capture:** Monitors the database's write-ahead log (WAL) or oplog, capturing every committed change as an event. This happens at the database level, not application level, guaranteeing we never miss changes.
3. **Immutable Event Stream:** CDC publishes all changes to Kafka, creating an append-only log of every state transition. Events are keyed by `payment_intent_id` and include the full before/after state.
4. **Specialized Consumers:** Different services consume the event stream for their specific needs without impacting the operational database.


```
# Example Kafka events from CDC
{
  "op": "insert",
  "source": "payment_intents_db",
  "table": "payment_intents",
  "ts_ms": 1681234567890,
  "before": null,
  "after": {
    "payment_intent_id": "pi_123",
    "merchant_id": "merch_456",
    "amount": 2500,
    "status": "created"
  }
}

{
  "op": "update",
  "source": "payment_intents_db",
  "table": "payment_intents",
  "ts_ms": 1681234568901,
  "before": {
    "payment_intent_id": "pi_123",
    "status": "created"
  },
  "after": {
    "payment_intent_id": "pi_123",
    "status": "authorized"
  }
}
```

The nice part about this architecture is that different consumers can materialize different views:

- **Audit Service:** Maintains a complete, immutable history optimized for compliance queries
- **Analytics:** Builds denormalized views for business intelligence
- **Reconciliation:** Correlates our events with external payment network events (covered in detail later)
- **Webhook Delivery:** Tracks which state changes need merchant notification (also covered later)

Granted, most of the above are out of scope for our requirements, but it's a useful callout that we recognize that the system will grow to require these things and this will make it easier.



For durability, Kafka provides configurable replication (typically 3x) across multiple brokers and availability zones.

Events are retained for a configurable period (often 7-30 days) on disk, with older events archived to object storage for permanent retention. All events are automatically flushed to S3 for long-term auditability, ensuring we can reconstruct any payment's complete history even years later for compliance audits or dispute resolution. This gives us both hot data for operational use and cold storage for compliance.

This architecture provides the best of both worlds: merchants get sub-10ms API responses from an optimized operational database, while every change is automatically captured in an immutable event stream without impacting API latency. Since CDC operates at the database level, there's no reliance on application code remembering to write audit records—if a change commits to the database, it will appear in the event stream.

The separation of concerns allows independent scaling: your operational database handles 10k TPS focused purely on current state, while the event stream and its consumers handle audit, analytics, and reconciliation without competing for resources. This also satisfies compliance requirements with cryptographically verifiable immutability and provides the foundation for advanced capabilities like reconciliation and webhook delivery.

By implementing this hybrid approach, we achieve true durability and auditability without sacrificing the performance our merchants expect. The immutable event stream becomes the foundation for not just audit compliance, but for the advanced capabilities—reconciliation, analytics, and real-time notifications—that distinguish a professional payment system from a simple transaction processor.



A savvy interviewer might ask: "Isn't CDC a single point of failure? What happens if your CDC system fails and you miss critical payment events?"

This is a great question! CDC is technically a single point of failure - if it stops working, events stop flowing to Kafka even though database writes continue. Companies like Stripe handle this by running multiple independent CDC instances reading from the same database, each writing to different Kafka clusters. They also implement monitoring that alerts within seconds if CDC lag increases, and maintain recovery procedures to replay missed events from database logs if needed. For critical payment events, they might also implement application-level fallbacks that write directly to Kafka if CDC hasn't confirmed the event within a certain timeframe.

3) The system should guarantee transaction safety and financial integrity despite the inherently asynchronous nature of external payment networks

Payment networks operate in a fundamentally different way than our internal systems. When we send a charge request to Visa, Mastercard, or a bank, we're crossing into systems we don't control. These networks process millions of transactions across global infrastructure, with their own retry mechanisms, queue delays, and batch processing windows. A payment we consider "timed out" might still be winding its way through authorization systems, while another might have succeeded instantly but lost its response packet on the way back to us.

This asynchronous reality creates serious risks. The most dangerous is double-charging. A customer clicks "pay" for their \$50 purchase, we timeout waiting for the bank's response, the merchant retries, and suddenly the customer sees two \$50 charges on their statement. Even if we eventually refund one, we've damaged trust and created support headaches. The opposite risk is equally damaging: a payment succeeds at the bank but we never know it, so we tell the merchant the payment failed. The merchant never ships the product, the customer's card gets charged anyway, and now we have an angry customer who paid for goods they'll never receive.

The challenge isn't eliminating this asynchronous behavior, which is impossible. Instead, we need to build systems

that acknowledge uncertainty and handle it gracefully. Fortunately, the CDC and event stream infrastructure we established for durability provides exactly the foundation we need to track payment attempts, handle timeouts intelligently, and maintain consistency even when networks misbehave.

Let's explore different approaches to handling this inherent uncertainty:

Bad Solution: Assume Timeouts Mean Failure

Approach

A bad option is to just treat network communication like internal service calls. We set a reasonable timeout (typically 30-60 seconds) for payment network requests. If we don't receive a response within that window, we assume the payment failed and update our database accordingly.

This straightforward approach lets us give merchants quick feedback. After 30 seconds, they know to either fulfill the order or ask the customer to try again. Our database maintains a clean state with every payment marked as either succeeded or failed, never stuck in limbo.

Challenges

This approach is a disaster waiting to happen, and for obvious reasons. Network timeouts tell us nothing about what actually happened with the payment. Consider this scenario that happens daily in production systems:

1. We send a \$200 charge request to the customer's bank
2. The bank approves and debits the customer's account
3. The response packet gets delayed or lost in network congestion
4. Our 30-second timeout triggers, we mark the payment as "failed"
5. The merchant displays "Payment failed, please try again"
6. The customer retries, creating a second \$200 charge
7. The customer now has \$400 in charges for a \$200 purchase

Yikes!

Good Solution: Pending States with Manual Reconciliation

Approach

A better approach acknowledges that timeouts create uncertainty, not failure. Instead of guessing, we track unclear outcomes explicitly by expanding our payment status to include a "pending_verification" state specifically for timeouts. We also create a separate table to track every attempt we make to charge a payment network, recording what we sent, when we sent it, and what reference ID the network assigned.

When a timeout occurs, we mark the payment as "pending_verification" rather than failed. A background service then periodically queries the payment network using our recorded reference ID to determine what actually happened with the payment. This verification process runs independently of the merchant-facing API, gradually resolving uncertain states without blocking new payment processing.

Importantly, to prevent duplicate charges, we implement idempotency using a unique database constraint on merchant ID and idempotency key. When merchants retry with the same idempotency key, they get the existing payment record instead of creating a duplicate charge.

Challenges

The manual reconciliation burden is significant. Our financial team must review pending payments daily, check them against bank files, and update statuses manually. During busy periods like Black Friday, thousands of payments might be stuck in pending states, overwhelming our reconciliation processes.

More fundamentally, this approach still treats reconciliation as an exception case rather than a core system behavior. We're bolting verification onto a system designed for synchronous responses, creating operational complexity without addressing the root cause. The database bears the full load of both operational queries and reconciliation lookups, leading to performance degradation during peak times when we can least afford it.

Great Solution: Event-Driven Safety with Reconciliation

Approach

The most robust approach leverages the CDC and event stream infrastructure we established for durability, treating asynchronous payment networks as first-class citizens in our architecture.


The key insight is to track our intentions before acting on them. Here's how the complete flow works:

1. **Record the attempt:** Before calling any payment network, we write an attempt record to our database with the network name, reference ID, and what we're trying to accomplish. This triggers a CDC event capturing our intention.
2. **Call the payment network:** We send the actual charge request with our configured timeout.
3. **Handle the response** (branching based on outcome):
 - **Success:** Update the attempt status to "succeeded" in the database, which triggers another CDC event with the successful outcome
 - **Timeout:** Update the attempt status to "timeout" in the database, triggering a CDC event that the reconciliation service will process
 - **Explicit failure:** Update the attempt status to "failed" with the failure reason
4. **Automated reconciliation:** A dedicated reconciliation service consumes timeout events and proactively queries the payment network using our recorded reference ID to determine what actually happened.

This creates a complete audit trail where every step is captured as an immutable event, whether the call succeeds, fails, or times out.

The most powerful aspect is how we handle reconciliation. Payment networks provide two ways to verify transaction status: direct API queries for real-time verification and batch reconciliation files. These reconciliation files are comprehensive records of all processed transactions during a specific period, delivered on regular schedules (daily or hourly). They follow strict formatting specifications and serve as the definitive record of what actually happened in the payment network.

A dedicated reconciliation service correlates these external events with our internal attempts, updating payment statuses based on authoritative network data rather than guessing from timeouts. The reconciliation service continuously consumes events, proactively queries networks when timeouts occur, and systematically processes settlement files when they arrive.



This architecture acknowledges that payment networks are asynchronous partners, not synchronous services. By tracking every attempt and its outcome through events, we build a system that handles uncertainty gracefully rather than fighting it.

! In production payment systems like Stripe, transaction processing actually uses a two-phase event model:

1. **Transaction Created Event:** Emitted when the transaction service begins processing, before the database write is complete. If this event fails to emit, the transaction enters a locked/failed state and can be retried.
2. **Transaction Completed Event:** Emitted after the database write has successfully completed. If this event fails to emit, the transaction also enters a locked/failed state where further updates are blocked until the completion event is successfully emitted.

This two-phase approach allows the system to compare the "created" event data with the actual database state during retries. If the database write already occurred, the system only needs to re-emit the completion event rather than retry the entire transaction. This pattern provides stronger guarantees about data consistency between the transaction service and downstream consumers.

Most missing "completed" events in production are actually due to external system timeouts before the database write occurred, rather than event emission failures after successful writes.

The key to handling asynchronous payment networks is accepting that uncertainty is inevitable and building systems designed for eventual consistency. By combining a traditional database for synchronous merchant needs with an event stream for asynchronous network reality, we achieve both performance and correctness. This pattern, proven at companies like Stripe processing billions in payments, shows that the best distributed systems don't fight the nature of external dependencies — they embrace and design for them.

4) The system should be scalable to handle high transaction volume (10,000+ TPS)

We've come a long way! We have a payment processing system that meets just about all of our functional and non-

functional requirements. Now we just need to discuss how we would handle scale. I like to save scale for my last deep dive because then I have a clear understanding of the full system I need to scale. It's likely that you already talked about scale (as we did a bit) during other parts of the interview, but we'll do our best to tie it all together here.

Servers

Let's start by acknowledging the basics. Each of our services is stateless and will scale horizontally with load balancers in front of them distributing the load. This is worth mentioning in the interview, but there is no need to draw it out; your interviewer knows what you mean, and this is largely taken for granted nowadays.

Kafka

Next, let's look deeper at our event log, Kafka. We are expecting ~10k tps. While Kafka can support millions of messages per second at the cluster level, it's important to note that each individual partition typically handles around 5,000-10,000 messages per second under normal production conditions. To comfortably support our 10k TPS requirement, we'd need multiple partitions—likely 3-5—to ensure both throughput and fault tolerance.

When partitioning, we need to guarantee ordering for transactions per PaymentIntent. This suggests partitioning our Kafka topics by payment_intent_id, which ensures all events for a given PaymentIntent are processed in order while allowing parallel processing across different PaymentIntents. This partitioning strategy balances our throughput needs with our ordering requirements—all state transitions for a single PaymentIntent (created → authorized → captured) will be processed sequentially, while different PaymentIntents can be processed in parallel across partitions.

For redundancy, we'd configure Kafka with a replication factor of 3, meaning each partition has three copies distributed across different brokers. This provides fault tolerance if a broker fails. We'd also set up consumer groups for each of our services, allowing us to scale consumers horizontally while maintaining exactly-once processing semantics.



Learn more about how to scale Kafka in our [Kafka deep dive here](#).

Database

Our database should be approximately the same order of magnitude as our event log with regards to the number of transactions per second. With 10k TPS for our event stream, we can expect around 10k writes per second to our database as well. This is right on the edge of what a well-optimized PostgreSQL instance can handle this load, especially when combined with read replicas and proper indexing. To better handle the scale, we can shard our database by merchant_id.

With regards to data growth, we can estimate that each row is ~500 bytes. This means we are storing $10,000 * 500 \text{ bytes} = 5\text{mb}$ of data a second, $5\text{mb} * 100,000 \text{ (rounded seconds in a day)} = 500\text{gb}$ of data a day, and $500\text{gb} * 365 = \sim 180\text{tb}$ of data a year. This is significant storage growth that will require careful planning for data retention and archiving strategies.

For transactions older than a certain period (e.g., 3-6 months), we can move them to cold storage like Amazon S3 or Google Cloud Storage. This archived data would still be accessible for compliance and audit purposes but wouldn't impact our operational database performance. We'd implement a scheduled job to identify, export, and purge old

records according to our retention policy.

For read scaling, we can implement read replicas of our database. Most payment queries are read operations (checking payment status, generating reports), so having multiple read replicas will distribute this load. We can also implement a caching layer using Redis or Memcached for frequently accessed data like recent payment statuses.

Bonus Deep Dives

1) How can we expand the design to support Webhooks?



As mentioned earlier, a robust webhook system is a substantial system design topic in its own right - one that could easily fill an entire interview (and often does!). In the context of this payment processor discussion, we'll outline the high-level approach to webhooks rather than diving into all implementation details.

While our polling-based status endpoint works well for basic scenarios, merchants often need real-time updates about payment status changes to trigger business processes like order fulfillment or access provisioning. Webhooks solve this by allowing our system to proactively notify merchants about events as they occur.


Merchants provide us with two additional bits of information:

1. Callback Url: This is the URL that we will **POST** updates to when we have them.
2. Subscribed Events: This is a list of events that they want to subscribe to. We will notify them, at the callback url, when any of these events occur.



You hear real-time updates and you may think websockets or SSE! But no, these are often conflated concepts, but it's important to understand that webhooks represent server-to-server communication, not server-to-client (like websockets or SSE). The payment system's server sends notifications directly to the merchant's server via HTTP requests to a predefined endpoint. This is fundamentally different from client-facing real-time updates (like WebSockets or Server-Sent Events), which would deliver updates from a server to a browser or mobile app. Webhooks are designed for system-to-system communication and typically require the merchant to operate a publicly accessible endpoint to receive these notifications.

Here's how webhooks would work at a high level in our payment processing system:



1. **Database Changes:** Our Transaction and PaymentIntent services update the operational database as payments progress through their lifecycle (created → authorized → captured, etc.).
2. **CDC Events:** Change Data Capture automatically captures these database changes and publishes them to our Kafka event stream. These CDC events include payment status changes, transaction completions, and other state transitions.
3. **Webhook Service:** We introduce a new Webhook Service that consumes from the same Kafka event stream as our other specialized consumers ie, Reconciliation. When the service receives a CDC event, it checks if the associated merchant has configured a webhook endpoint for that event type. If configured, it prepares the webhook payload with relevant event details, signs the payload with a shared secret to enable verification, and attempts delivery to the merchant's endpoint.
4. **Delivery Management:** For each webhook, the delivery attempt is recorded with its status. If delivery fails, the Webhook Service implements a retry strategy with exponential backoff (e.g., retry after 5s, 25s, 125s, etc., up to a reasonable maximum interval like 1 hour).
5. **Merchant Implementation:** On the merchant side, they would need to configure a publicly accessible HTTPS endpoint to receive webhooks. They must verify the signature of incoming webhooks using the shared secret to ensure authenticity. After verification, they would process the webhook payload and update their systems accordingly with the new information. Finally, they should return a 2xx HTTP status code to acknowledge receipt of the webhook, preventing unnecessary retries.

Simple example of a webhook payload:

```
{
  "id": "evt_1JklMnOpQrStUv",
  "type": "payment.succeeded",
  "created": 1633031234,
  "data": {
    "object": {
      "id": "pay_1AbCdEfGhIjKlM",
      "amountInCents": 2499,
      "currency": "usd",
      "status": "succeeded",
      "created": 1633031200
    }
  }
}
```



In production systems, webhook delivery reliability is critical. Although we're not diving into all the details here, a complete webhook system would include features like idempotency keys, detailed delivery monitoring, webhook logs for debugging, and a dashboard for merchants to view and replay webhooks.

If this were a dedicated webhook system design interview, we would explore more intricate challenges such as exactly-once delivery semantics, handling webhook queue backlogs during outages, webhook payload versioning, and adaptive rate limiting to prevent overwhelming merchant systems. However, for our payment processor design, this high-level overview captures the essential functionality that would complement our core payment flows.

What is Expected at Each Level?

You may be thinking, "how much of that is actually required from me in an interview?" Let's break it down.

Mid-level

For a mid-level candidate discussing a payment system, I'd primarily focus on your ability to establish the core functionality. You should be able to design a straightforward payment flow that handles the basic requirements of initiating payment requests, processing cards, and returning status updates.

You should demonstrate understanding of the need for security in payment systems and identify that sensitive card data shouldn't touch merchant servers. However, I wouldn't expect you to deeply understand all the nuances of payment security or complex tokenization approaches. When discussing consistency, a reasonable solution like assuming timeouts mean failure (even if not ideal) would be acceptable as long as you can explain the approach clearly. I'd be looking more at your problem-solving process than expecting a perfect solution to complex payment reconciliation challenges.

I'd anticipate that you might need guidance on deeper technical challenges, and that's perfectly fine. The interviewer will likely direct the conversation through the latter portions of the discussion.

Senior

For a senior engineer, I'd expect you to move quickly through the basic functionality and drive the conversation toward the critical non-functional aspects of a payment system. You should proactively identify security and consistency as core challenges.

On the security front, you should propose a solution like the `iframe` isolation approach and understand its benefits in protecting sensitive card data. You should be able to articulate why this approach is superior to having merchants collect card data directly. For payment consistency, you should recognize that timeouts don't necessarily mean failure and propose a more sophisticated approach like `idempotent` transactions with pending resolution. You should be able to identify potential race conditions and explain how your design prevents double-charging.

For scaling discussions, I'd expect you to speak with confidence about how you'd implement horizontal scaling for services and databases, though you might need some prompting to explore more advanced topics like Kafka partitioning strategies or event sourcing.

Staff+

As a staff+ candidate, I'd expect you to demonstrate deep expertise in payment systems design while also showing exceptional product and architectural thinking. Rather than rushing to complex solutions, you should first evaluate whether they're necessary. For instance, you might point out that for many payment scenarios, event sourcing with reconciliation is the right approach not just for technical reasons but because it aligns with how financial systems fundamentally work - acknowledging the inherently asynchronous nature of payment networks. For security, you should be able to explain multi-layered approaches like tokenization with client-side encryption, showing an understanding of the defense-in-depth strategy necessary for handling sensitive financial data.

You should proactively identify and address the most challenging edge cases in payment processing. For example,

you might discuss how to handle payment network outages through fallback processing paths or how to design resilient reconciliation processes that guarantee eventual consistency with external payment networks.



Test Your Knowledge

Take a quick 15 question quiz to test what you've learned.

Start Quiz

Mark as read ☐

[Next: Real-time Updates](#) →

How would you rate the quality of this article?

1 2 3 4 5 6 7 8 9 10

Poor

Great

Add a comment...

B I </> 99 ≡ ≡ ≡ ≡

Comment

☐ Anonymous

Posting as leetcode

Search 211 comments

Sort By

Popular



Siddharth Chaudhary ★ Top 5% • 7 months ago

I feel like designing a payment system requires so much domain knowledge!

I think it's unfair to ask this question in a system design interview because the concepts you learned from designing other products/services are not applicable here.

What do you think @Evan?

42 [Reply](#)



Evan King Admin • 7 months ago

Completely agree. I was a bit surprised to see it requested as much as it was. It's not a question I would ask without providing a ton of background knowledge to the candidate upfront.

24 [Reply](#)



Siddharth Chaudhary ★ Top 5% • 7 months ago

Yeah, that's because some companies ask this question to raise the difficulty level of their interviews especially in India. And everyone was asking for it because it's hard to find a good source to study this from. Thanks so much for taking the time to break this problem down. Learnt some new things, very helpful!!

16 [Reply](#)



Phuong Le Premium • 7 months ago

I guess this question is usually asked by fintech companies e.g. checkout.com because it is closed to their internal systems, which helps them discuss it more comfortably, and also it can tell if a candidate have the relevant domain knowledge, which is a great bonus for the candidate.

3 [Reply](#)

 [reply](#)

YelpingAmaranthBaboon850 Premium • 5 months ago

I was asked almost this question when I interviewed at Stripe last week.

 3 [Reply](#)



Evan King Admin • 5 months ago

How'd it go? 🍌

 1 [Reply](#)

YelpingAmaranthBaboon850 Premium • 5 months ago

Poorly. I hadn't seen this guide; just got a notification about it today. I was underprepared and the interviewer got really agitated about my API design immediately and wouldn't let me iterate on it or proceed. It was a weird experience.

 5 [Reply](#)



Evan King Admin • 5 months ago

Dang, super sorry to hear that 😞 Sounds like a frustrating experience

 4 [Reply](#)

M **MatureOliveHalibut263** Premium • 7 months ago

Hey Evan, thanks for your content. I got asked the dropbox question during a mock interview and was asked lots of questions that was not covered in your design (not saying it should). So I have been using ChatGPT to generate extra questions based on your solutions, but its not as good.

I'd strongly recommend adding 5-10 questions at the end where users can brainstorm themselves. You don't have to provide answers. For example, one question I was asked on dropbox question was: what if there are millions of small files and you need to upload them all. Apparently dropbox faced this problem in production and had to add a feature to compress and bulk upload small files to save on time.

Thanks again, love ur content!

 16 [Reply](#)

F **FlyingTomatoHeron745** Premium • 7 months ago

What follow-up questions were you asked?

 0 [Reply](#)

K **Karthik555** Premium • 6 months ago

I think you can explain the concept of chunking its and provide and track it from the backend, customized for Dropbox.

 0 [Reply](#)

G **ghanekar.omkar** Premium • 7 months ago

I'm not sure I understood the math (1K writes/sec) under the Database section of 3rd deep dive.

<https://www.hellointerview.com/learn/system-design/problem-breakdowns/payment-system#database>

Our database, given that it is storing projections of data, should be up to an order of magnitude smaller than our event log with regards to the number of transactions per second. 1k write per second in the context of modern databases is not a lot, we can handle that with a single well optimized postgres instance.

Our initial non fn requirement states 10K TPS which I assume to mean 10K unique payment transactions and not 10K overall interim states.

If that assumption is right, it'd mean 10K DB entries/second and many more messages on the queue due to each transaction going through up to 4-5 interim states.

Is my assumption right or am I misinterpreting something?

👍 10 [Reply](#)

A

Abhi ★ Top 10% • 6 months ago

I have the same question!

👍 0 [Reply](#)



Mikhail Tcirulnikov Premium • 8 months ago

Thank you for event sourcing approach description here. It is brilliant application for it. How long we want to store events in Kafka?

👍 8 [Reply](#)



Evan King Admin • 8 months ago

Great question, you typically want to retain events in Kafka for ~30 days max. The event log isn't meant for long-term storage, you'd archive the events to cold storage (S3, GCS) for audit/compliance, and maintain derived state in your databases.

👍 23 [Reply](#)

W

walnatar2 Premium • 6 months ago

I think you need have regular snapshot/checkpointing also, so we can reply from derived state

👍 0 [Reply](#)

U

UltimateBronzeLemur440 • 7 months ago

The description of the client side encryption approach is incorrect and misses the point of PCI. Even encrypted card numbers are considered pci data, so passing an encrypted card number to your server would bring that entire server, api gateway, network, and all connected services into PCI, which is not ideal. The point of the iframe is to load a form that comes from a pci compliant domain that is different from the parent origin.

Obviously only people who have worked on payment systems would know this in detail, so that brings me to my question: what is the point of this site or even these interviews? Is it for people to memorize the answers to commonly asked questions (cheat), or is it to show problem solving skills. I'm genuinely curious (and only being a little combative).

👍 7 [Reply](#)



Evan King Admin • 7 months ago

The goal definitely isn't to memorize perfect answers (you won't know every domain and perfect answers don't even exist), but to show you can reason about tradeoffs and constraints while being open to guidance when discussing unfamiliar territory. A payment system interview isn't testing if you know PCI compliance details, it's testing if you recognize that handling sensitive financial data requires special consideration and can propose reasonable approaches to address that. Or, given input from your interviewer, can you reason about the tradeoffs of multiple solutions. Like LeetCode, system design is about patterns. These articles aim to give candidates exposure to patterns they can apply to novel problems, not "correct" answers to memorize.

👍 37 [Reply](#)

[Show All Comments](#)

Schedule a mock interview

Meet with a FAANG senior+ engineer or manager and learn exactly what it takes to get the job.

[Schedule A Mock Interview](#)

Questions

[Meta SWE Interview Questions](#)
[Amazon SWE Interview Questions](#)
[Google SWE Interview Questions](#)
[OpenAI SWE Interview Questions](#)
[Engineering Manager \(EM\) Interview Questions](#)

Learn

[Learn System Design](#)
[Learn DSA](#)
[Learn Behavioral](#)
[Learn ML System Design](#)
[Learn Low Level Design](#)
[Guided Practice](#)

Links

[FAQ](#)
[Pricing](#)
[Gift Mock Interviews](#)
[Gift Premium](#)
[Become a Coach](#)
[Our Coaches](#)
[Hello Interview Premium](#)

Legal

[Terms and Conditions](#)
[Privacy Policy](#)

Contact

[About Us](#)
[Product Support](#)
7511 Greenwood Ave North
Unit #4238 Seattle
WA 98103